

Chapter 1 – Conditioning and stability

IEEE 754 double-precision format Normalized numbers are written

$$x = \pm(1.b_1 \dots b_p)_2 \times 2^{(E-1023)},$$

with precision $p = 52$. The significand has 52 bits (with hidden leading 1), the exponent 11 bits with $1 \leq E \leq 2046$, and 1 sign bit.

Range The smallest and largest normalized numbers are

$$x_{\min} = 2^{-1022} \approx 2.2 \cdot 10^{-308}, \quad x_{\max} = (2-2^{-52}) \times 2^{1023} \approx 1.8 \cdot 10^{308}.$$

Machine precision Gap between 1 and next floating point number,

$$\varepsilon_M = 2^{-(p)} = 2^{-52} \approx 2.2 \cdot 10^{-16}.$$

Subnormal numbers These are of the form

$$x = \pm(0.b_1 \dots b_{p-1})_2 \times 2^{E_{\min}-1023}.$$

These fill the gap below $x_{\min}$, smallest is $2^{-1074} \approx 5 \cdot 10^{-324}$.

Exceptions An exponent $E = 0$ with zero significand encodes ± 0 , while $E = 2047$ encodes $\pm \infty$ (if significand=0) or NaN (otherwise).

Rounding (Thm. 1.2) For any x such that $x_{\min} < |x| < x_{\max}$,

$$\hat{x} = x(1 + \varepsilon), \quad |\varepsilon| \leq \varepsilon_M/2 \quad (\text{nearest}).$$

Floating point arithmetic (Thm. 1.4) Let x and y be two normalized floating point numbers such that $x_{\min} < |x \star y| < x_{\max}$. Then

$$\widehat{x \star y} = (x \star y)(1 + \varepsilon), \quad \star \in \{+, -, \times, /\}, \quad |\varepsilon| \leq \varepsilon_M/2 \quad (\text{nearest}).$$

Relative condition number (Def. 1.9) It is defined by

$$\kappa(x) = \lim_{\varepsilon \rightarrow 0} \sup_{\|\delta x\| \leq \varepsilon} \left(\frac{\|\delta f\|}{\|f(x)\|} \right) \bigg/ \frac{\|\delta x\|}{\|x\|}.$$

For differentiable f with Jacobian J , $\kappa(x) = \|J(x)\| / (\|f(x)\| \|x\|)$.

Forward stability (Def. 1.10) A stable algorithm gives nearly the right answer to nearly the right question, i.e.,

$$\frac{\|\hat{f}(x) - f(\hat{x})\|}{\|\hat{f}(\hat{x})\|} = \mathcal{O}(\varepsilon_M) \quad \text{with} \quad \frac{\|\hat{x} - x\|}{\|x\|} = \mathcal{O}(\varepsilon_M).$$

Backward stability (Def. 1.11) A backward stable algorithm gives exactly the right answer to nearly the right question, i.e.,

$$\hat{f}(x) = f(\hat{x}) \quad \text{with} \quad \frac{\|\hat{x} - x\|}{\|x\|} = \mathcal{O}(\varepsilon_M).$$

These verify $\|\hat{f}(x) - f(\hat{x})\| / \|f(\hat{x})\| = \mathcal{O}(\kappa(x)\varepsilon_M)$ (Prop. 1.13).

Chapter 2 – Polynomial interpolation

Uniqueness (Thm. 2.1) Given $(n+1)$ distinct $x_0, \dots, x_n \in \mathbb{R}$ and $(n+1)$ values $y_0, \dots, y_n \in \mathbb{R}$, there is a unique polynomial p_n of degree at most n such that $p_n(x_i) = y_i$ for all $i = 0, \dots, n$.

Lagrange polynomials These are defined by

$$\ell_i(x) = \prod_{j \neq i} \frac{x - x_j}{x_i - x_j}$$

and verify $\ell_i(x_j) = \delta_{ij}$. We also define $\Pi_{n+1}(x) = \prod_{i=0}^n (x - x_i)$.

Errors for differentiable functions (Thm. 2.2) Assume $f : [a, b] \rightarrow \mathbb{R}$ is $\mathcal{C}^{n+1}$. Then for any set of distinct points $x_0, \dots, x_n$,

$$\|f - p_n(f)\|_{\infty} \leq \frac{\|\Pi_{n+1}\|_{\infty}}{(n+1)!} \|f^{(n+1)}\|_{\infty}.$$

Errors for analytic functions (Thm. 2.3) Assume $f : [a, b] \rightarrow \mathbb{R}$ is analytic in the disk centered at $(a+b)/2$ with radius $R > 3/2(b-a)$. Then for any set of distinct points $x_0, \dots, x_n$,

$$\|f - p_n(f)\|_{\infty} \rightarrow 0 \quad \text{as } n \rightarrow \infty.$$

Lebesgue cst (Def. 2.10) Defined by $\Lambda_n = \sup_{\|f\|_{\infty} \leq 1} \|p_n(f)\|_{\infty}$.

Error for Lipschitz functions (Thm. 2.16) Assume $f : [a, b] \rightarrow \mathbb{R}$ is Lipschitz. Then for any set of distinct points $x_0, \dots, x_n$,

$$\|f - p_n(f)\|_{\infty} \leq C \Lambda_n / \sqrt{n}.$$

For Chebyshev points, $\Lambda_n \sim \log(n)$, hence $\|f - p_n(f)\|_{\infty} \rightarrow 0$.

Errors for differentiable functions (Cor. 2.18) Assume $f : [a, b] \rightarrow \mathbb{R}$ is $\mathcal{C}^k$. Then for Chebyshev points,

$$\|f - p_n(f)\|_{\infty} \leq C_k \frac{\log n}{n^k}.$$

Errors for analytic functions (Thm. 2.19) Assume $f : [a, b] \rightarrow \mathbb{R}$ is analytic in the Bernstein ellipse E_{ρ} . Then for Chebyshev points,

$$\|f - p_n(f)\|_{\infty} \leq C \frac{\sup_{z \in E_{\rho}} |f(z)|}{\rho^n}.$$

Chapter 3 – Quadrature

Elementary quadrature rule (Def. 3.1) An $(n+1)$ -point elementary quadrature rule is a linear form $I_n : \mathcal{C}^0([a, b]) \rightarrow \mathbb{R}$ of the form

$$I_n(f) = \sum_{k=0}^n w_k f(x_k).$$

Lagrange-based quadrature (Def. 3.3) When $w_k = \int_a^b \ell_k(x) dx$. Equidistant/Chebyshev gives Newton–Cotes/Clenshaw–Curtis.

Errors for differentiable functions (Thm. 3.18) For an $(n+1)$ -point Lagrange-based rule and $f \in \mathcal{C}^{n+1}([a, b])$,

$$|I(f) - I_n(f)| \leq \frac{\|f^{(n+1)}\|_{\infty}}{(n+1)!} \int_a^b |\Pi_{n+1}(x)| dx.$$

Error for analytic functions (Thm. 3.19) Assume $f : [a, b] \rightarrow \mathbb{R}$ is analytic in the disk centered at $(a+b)/2$ with radius $R > 3/2(b-a)$. For any sequence of Lagrange-based quadrature rules,

$$|I(f) - I_n(f)| \rightarrow 0 \quad \text{as } n \rightarrow \infty.$$

Errors for differentiable functions (Thm. 3.25) If f is in $\mathcal{C}^k([a, b])$:

$$|I(f) - I_n(f)| \leq \begin{cases} C/n^k, & \text{Clenshaw–Curtis,} \\ C/(2n+1)^k, & \text{Gauss–Legendre.} \end{cases}$$

Errors for analytic functions (Thm. 3.26) If f is analytic in E_{ρ} :

$$|I(f) - I_n(f)| \leq \begin{cases} CM_{\rho}/\rho^n, & \text{Clenshaw–Curtis,} \\ CM_{\rho}/\rho^{2n}, & \text{Gauss–Legendre.} \end{cases}$$

Composite quadrature (Def. 3.27) Given a subdivision $a = t_0 < t_1 < \dots < t_m = b$, a composite quadrature rule is the linear form

$$I_{m,n}(f) = \sum_{j=0}^{m-1} I_n^{[t_j, t_{j+1}]}(f), \quad I_n^{[t_j, t_{j+1}]} \text{ elementary rule.}$$

Trapezoidal rule (Ex. 3.30) The composite trapezoidal rule is

$$I_m(f) = \frac{b-a}{m} \sum_{k=0}^{m-1} \frac{f(x_k) + f(x_{k+1})}{2}, \quad x_k = a + k(b-a)/m.$$

For $f \in \mathcal{C}^2([a, b])$, $I_m(f) = I(f) + \mathcal{O}(m^{-2})$. This a composite quadrature rule with Newton–Cotes ($n = 1$) on each subinterval.

Chapter 4 – Direct methods for linear systems

Gaussian elimination (Thm. 4.1) For any $A \in \mathbb{R}^{n \times n}$, there exists a permutation matrix P such that

$$PA = LU \quad (A \in \mathbb{R}^{n \times n}),$$

where L is a unit lower triangular matrix ($\ell_{ii} = 1$) and U is upper triangular. If A is nonsingular, U is nonsingular and the factorization can be used to solve $Ax = b$. Computational cost is $\mathcal{O}(2n^3/3)$.

Condition number κ (Thm. 4.3) Let $A \in \mathbb{R}^{n \times n}$ be nonsingular, and consider the systems $Ax = b$ and $\hat{A}\hat{x} = \hat{b}$ with

$$\|\hat{A} - A\| \leq \varepsilon_A \|A\|, \quad \|\hat{b} - b\| \leq \varepsilon_b \|b\|, \quad \varepsilon_A, \varepsilon_b \leq \varepsilon_M.$$

If $\varepsilon_A \kappa(A) < 1$, then

$$\frac{\|\hat{x} - x\|}{\|x\|} \leq \frac{\kappa(A)}{1 - \varepsilon_A \kappa(A)} (\varepsilon_A + \varepsilon_b),$$

where $\kappa(A) = \|A\| \|A^{-1}\|$ is the condition number of A .

Properties of κ (Prop. 4.4) Let $A \in \mathbb{R}^{n \times n}$ be nonsingular. Then

- (a) $\kappa(A) \geq 1$;
- (b) $\kappa(\alpha A) = \kappa(A)$ for any $\alpha \neq 0$;
- (c) $\kappa(A) = \frac{\max_{\|y\|=1} \|Ay\|}{\min_{\|z\|=1} \|Az\|}$.

Property (c) allows extending the definition to rectangular matrices.

Stability of Gaussian elimination (Thm. 4.5) Let $A \in \mathbb{R}^{n \times n}$ be nonsingular, and let $\hat{A} = \hat{L}\hat{U}$ be computed by Gaussian elimination with complete pivoting (so that $|\ell_{ij}| \leq 1$). Then

$$|\hat{a}_{ij} - a_{ij}| \leq 2a \min(i-1, j) \varepsilon_M, \quad a = \max_{i,j,k} |a_{ij}^{(k)}|.$$

Cholesky factorization (Thm. 4.6) Let $A \in \mathbb{R}^{n \times n}$ be symmetric positive definite (SPD). Gaussian elimination can be performed without pivoting and yields the Cholesky factorization

$$A = LL^T \quad (\text{SPD } A \in \mathbb{R}^{n \times n}),$$

where L is lower triangular. Computational cost is $\mathcal{O}(n^3/3)$.

Stability of Cholesky factorization (Thm. 4.7) Let $A \in \mathbb{R}^{n \times n}$ be SPD, and let $\hat{A} = \hat{L}\hat{L}^T$ be its Cholesky factorization. Then

$$|\hat{a}_{ij} - a_{ij}| \leq a_0 \min(i, j) \varepsilon_M, \quad a_0 = \max_{i,j} |a_{ij}| = \max_i |a_{ii}|.$$

Normal equations (Thm. 4.8) $A \in \mathbb{R}^{m \times n}$, $b \in \mathbb{R}^m$, $m > n$. Then

$$\|Ax - b\|_2 \rightarrow \min \iff A^T Ax = A^T b \quad (\text{normal equations}).$$

The normal equations always have at least one solution. Moreover, $A^T A$ is symmetric and non-negative since $x^T A^T A x = \|Ax\|_2^2 \geq 0$.

QR Factorization For any $A \in \mathbb{R}^{m \times n}$ with $m \geq n$, there exists an orthogonal matrix Q and an upper triangular matrix R such that

$$A = QR \quad (A \in \mathbb{R}^{m \times n}).$$

For least squares problems, only the first n columns of Q are needed. Computational cost is $\mathcal{O}(2mn^2 - 2n^3/3)$.

Solving the normal equations

1. Cholesky factorization: Compute $A^T A = LL^T$ and solve

$$x = L^{-T} L^{-1} A^T b.$$

This is faster, $\mathcal{O}(mn^2 + n^3/3)$ flops, but squares the condition number $\kappa(A^T A) = \kappa(A)^2$, hence less stable.

2. QR factorization: Compute $A = QR$ and solve

$$x = R^{-1} Q^T b.$$

This is slower, $\mathcal{O}(2mn^2 - 2n^3/3)$ flops, but more stable.

Chapter 5 – Iterative methods for linear systems

Direct vs iterative For sparse PDE systems, direct methods are efficient in 1D and 2D. In 3D, however, fill-in makes LU/Cholesky factorization costly in both memory and computation, so iterative methods with preconditioning are generally preferred.

Principle Consider $Ax = b$ in $\mathbb{R}^n$. Iterative methods generate a sequence $(x^k)_{k \in \mathbb{N}}$ converging to the solution x^∞ of $Ax = b$, typically via a fixed-point scheme $x^{k+1} = f_k(x^k)$.

- Stationary: x^{k+1} is a fixed combination of x^k and b , e.g.,

$$Mx^{k+1} = Nx^k + b \quad (A = M - N \text{ with cheap } M^{-1}).$$

- Non-stationary: the update $f_{k+1} \neq f_k$ varies at each iteration.

Convergence (Cor. 5.11) The iterative sequence

$$x^{k+1} = Px^k + c$$

converges for any $x^0 \in \mathbb{R}^n$ if and only if $\rho(P) < 1$. In that case

$$\lim_{k \rightarrow \infty} x^k = (I - P)^{-1} c.$$

Remark Substituting $P = M^{-1}N$ and $c = M^{-1}b$ gives

$$x^\infty = (I - M^{-1}N)^{-1} M^{-1}b.$$

For $A = M - N$, we have $M(I - M^{-1}N) = A$, hence $x^\infty = A^{-1}b$.

Gershgorin disks (Thm. 5.14) Let $A \in \mathbb{R}^{n \times n}$, and λ an eigenvalue of A . Then there exists $1 \leq i \leq n$ such that $|\lambda - a_{ii}| \leq \sum_{j \neq i} |a_{ij}|$.

Jacobi Choose $M = D$ (diagonal of A) and $N = D - A$:

$$x^{k+1} = (I - D^{-1}A)x^k + D^{-1}b.$$

Convergence: $\rho(I - D^{-1}A) < 1$ (Cor. 5.25); true if A is SRDD.

Gauss–Seidel Choose $M = L + D$ (lower triangular part plus diagonal) and $N = L + D - A = -U$ (strict upper triangular part):

$$x^{k+1} = -(L + D)^{-1} U x^k + (L + D)^{-1} b.$$

Convergence: $\rho((L + D)^{-1} U) < 1$; true if A is SRDD or SPD.

Spectral radius For $A \in \mathbb{R}^{n \times n}$, $\rho(A) = \max\{|\lambda| : \lambda \in \text{Sp}(A)\}$.

Spectral radius properties

- $\lim_{k \rightarrow \infty} \|A^k\|^{1/k} = \rho(A)$;
- $\rho(A) = \rho(S^{-1}AS)$ for any nonsingular S .

Matrix norms Let $A \in \mathbb{R}^{m \times n}$. Standard norms:

- 1-norm: $\|A\|_1 = \max_{1 \leq j \leq n} \sum_{i=1}^m |a_{ij}|$ (max column sum);
- inf-norm: $\|A\|_\infty = \max_{1 \leq i \leq m} \sum_{j=1}^n |a_{ij}|$ (max row sum);
- Frobenius norm: $\|A\|_F = \sqrt{\sum_{i,j} |a_{ij}|^2}$;
- Spectral norm (2-norm): $\|A\|_2 = \sqrt{\rho(A^T A)}$.

Matrix norms properties

- $\|Ax\| \leq \|A\| \|x\|$;
- $\|AB\| \leq \|A\| \|B\|$;
- $\|A\|_2 = \rho(A)$ if A is symmetric;
- $\|QA\|_2 = \|A\|_2$ and $\|QA\|_F = \|A\|_F$ for any orthogonal Q .

Chapter 6 – Solving nonlinear equations

Setup in $\mathbb{R}$ Let $f : \mathbb{R} \rightarrow \mathbb{R}$ be $\mathcal{C}^1$. We want to find a numerical solution to $f(x) = 0$. The true solution will be denoted by x_* .

Newton's method Given $x_0 \in \mathbb{R}$, the iterates are defined by

$$x_{n+1} = x_n - \frac{f(x_n)}{f'(x_n)}, \quad n \geq 0.$$

Secant method Given $x_0, x_1 \in \mathbb{R}$, the iterates are defined by

$$x_{n+1} = x_n - \frac{f(x_n)}{g(x_n)}, \quad g(x_n) = \frac{f(x_n) - f(x_{n-1})}{x_n - x_{n-1}}, \quad n \geq 1.$$

Convergence via contraction (Thm. 6.3) For $x_{n+1} = g(x_n)$ with fixed point $g(x_*) = x_*$, convergence holds where $|g'(x)| < 1$.

Convergence order (Def. 6.4) It is said to be of order $m \geq 1$ if

$$|x_{n+1} - x_*| = \mathcal{O}(|x_n - x_*|^m).$$

Newton convergence (Thm. 6.5) Let $f : [a, b] \rightarrow \mathbb{R}$ be $\mathcal{C}^2$ with a simple root x_* ($f(x_*) = 0$, $f'(x_*) \neq 0$). Then there exists $r > 0$ such that, for any $x_0 \in [x_* - r, x_* + r]$, Newton's iterates are well defined, stay in the interval, and converge quadratically to x_* .

Secant convergence (Thm. 6.7) Under the same assumptions, secant iterates converge superlinearly to x_* .

Stopping criteria The error $|x_{n+1} - x_*|$ is usually unavailable; practical criteria are the residual $|f(x_{n+1})|$ or the increment $|x_{n+1} - x_n|$.

Conditioning If f depends on data d , e.g., $d = (A, b)$ for linear systems, then the condition numbers can be approximated by

$$\kappa(d) \approx \frac{|d|}{|x_*| |f'(x_*)|}, \quad \kappa_{\text{abs}}(d) \approx \frac{1}{|f'(x_*)|}.$$

This yields

$$\frac{|\hat{x} - x_*|}{|x_*|} \lesssim \frac{|f(\hat{x})|}{|x_*| |f'(x_*)|}.$$

If well-conditioned, the residual $|f(\hat{x})|$ thus estimates the error well.

Setup in $\mathbb{R}^n$ Let $F : \mathbb{R}^n \rightarrow \mathbb{R}^n$ be $\mathcal{C}^1$. We want to find a numerical solution to $F(x) = 0$. The true solution will be denoted by x_* .

Newton's method Given $x_0 \in \mathbb{R}^n$, the iterates are defined by

$$x_{n+1} = x_n - J_F(x_n)^{-1} F(x_n), \quad n \geq 0.$$

Convergence order It is said to be of order $m \geq 1$ if

$$\|x_{n+1} - x_*\| = \mathcal{O}(\|x_n - x_*\|^m).$$

Newton convergence (Thm. 6.10) Let $F : \mathbb{R}^n \rightarrow \mathbb{R}^n$ be $\mathcal{C}^1$ on a convex open set D containing x_* . Assume $J_F(x_*)^{-1}$ exists with $\|J_F(x_*)^{-1}\| \leq C$ and J_F is L -Lipschitz in $B(x_*, R)$. Then there exists $r > 0$ such that, for any $x_0 \in B(x_*, r)$, Newton's iterates are well defined, stay in the ball, and converge quadratically to x_* .

Modified Newton's method To speed up Newton's method:

- freeze J_F for several iterations and reuse its LU factorization;
- solve inexactly with the exact J_F (e.g., a few Jacobi/SOR iterations yield Newton–Jacobi/Newton–SOR);
- solve exactly with an inexact J_F (e.g., finite differences).

Quasi-Newton methods Quasi-Newton methods approximate the Jacobian via updates $G_n \approx J_F(x_n)$. Let $\delta F_n = F(x_n) - F(x_{n-1})$ and $\delta x_n = x_n - x_{n-1}$. For instance, Broyden's method reads

$$x_{n+1} = x_n - G_n^{-1} F(x_n), \quad G_n = G_{n-1} + \frac{\delta F_n - G_{n-1} \delta x_n}{\|\delta x_n\|^2} \delta x_n^T.$$

Often, we initialize with $G_0 = J_F(x_0)$. Convergence holds under the same assumptions as Newton's method (Prop. 6.13).

Chapter 7 – ODEs I

Setup Consider, on $[t_0, t_0 + T]$, $y' = f(t, y)$ with $y(0) = y_0$.

Cauchy–Lipschitz (Thm. 7.5) Let f be continuous in t and uniformly Lipschitz continuous with respect to y for all $t \in [t_0, t_0 + T]$. Then there exists a unique differentiable solution y .

Time discretization Time is discretized by $N + 1$ points

$$t_0 < t_1 < \dots < t_N = t_0 + T,$$

with $\Delta t_n = t_{n+1} - t_n$, $\Delta t = \max \Delta t_n$. For $t_0 = 0$ and $\Delta t_n = \Delta t$:

$$t_0 = 0 < \dots < t_n = n\Delta t < \dots < t_N = T, \quad \Delta t = T/N.$$

General one-step methods These are of the form

$$y_{n+1} = y_n + \Phi(t_n, y_n; \Delta t_n).$$

For instance (see Ex. 8.2 when Φ is defined implicitly):

$$y_{n+1} = y_n + \Delta t_n f(t_n, y_n), \quad (\text{forward Euler}),$$

$$y_{n+1} = y_n + \Delta t_n f(t_{n+1}, y_{n+1}), \quad (\text{backward Euler}).$$

Errors & convergence (Def. 7.7) The global error is

$$e_n = y(t_n) - y_n, \quad n = 0, \dots, N.$$

The method is convergent if

$$\lim_{\Delta t \rightarrow 0} \left(\max_{n=0, \dots, N} |e_n| \right) = 0.$$

The local truncation error is defined by

$$\varepsilon_n = y(t_{n+1}) - y(t_n) - \Phi(t_n, y(t_n); \Delta t_n), \quad n = 0, \dots, N-1.$$

The consistency error is defined by

$$E = \sum_{n=0}^{N-1} |\varepsilon_n|.$$

Consistency (Def. 7.9) The method is consistent if

$$|\varepsilon_n| \leq \Delta t \Psi(\Delta t), \quad \Psi(\Delta t) \rightarrow 0 \quad \text{as} \quad \Delta t \rightarrow 0.$$

Convergence order (Def. 7.10) The scheme is of order $p \geq 1$ if

$$|\varepsilon_n| = \mathcal{O}(\Delta t_n^{p+1}), \quad n = 0, \dots, N-1.$$

Equivalently, if $E = \mathcal{O}(\Delta t^p)$.

Stability to perturbations (Def. 7.12) If there exists $C > 0$, independent of Δt , such that for any perturbed iteration

$$\xi_{n+1} = \xi_n + \Phi(t_n, \xi_n; \Delta t_n) + \eta_n,$$

it holds that

$$\max_{n=0, \dots, N} |y_n - \xi_n| \leq C \left(|y_0 - \xi_0| + \sum_{n=0}^{N-1} |\eta_n| \right).$$

This is equivalent to being zero-stable (bounded $\{y_n\}$ for $f = 0$).

Forward Euler We assume that f is L -Lipschitz.

1. Consistent, of order 1 if y is $\mathcal{C}^2$ (Prop. 7.11).
2. Stable with $C = e^{LT}$ (Prop. 7.14).
3. Convergent (Thm. 7.15), order 1 if y is $\mathcal{C}^2$ (Thm. 7.15).

Backward Euler We assume that f is L -Lipschitz.

1. Consistent if $L\Delta t < 1$, of order 1 if y is $\mathcal{C}^2$ (Prop. 7.27).
2. Stable with

$$C = \frac{e^{LT/(1-L\Delta t)}}{1-L\Delta t} \quad (\text{Prop. 7.26}).$$

3. Convergent (Thm. 7.28), order 1 if y is $\mathcal{C}^2$ (Thm. 7.29).

Absolute stability Abs. stable for $\lambda\Delta t$ if bounded $\{y_n\}$ for $f = \lambda y$.

Chapter 8 – ODEs II

Consistency/Order/Stability Consistency (Def. 8.3), convergence order (Def. 8.4), and stability (Def. 8.8) are defined as previously.

RK methods These are of the form $y_{n+1} = y_n + \Phi(t_n, y_n; \Delta t_n)$.

- RK1 (Euler method):

$$y_{n+1} = y_n + \Delta t_n f(t_n, y_n).$$

- RK2 (midpoint method):

$$\begin{aligned} k_1 &= f(t_n, y_n), \\ k_2 &= f\left(t_n + \frac{\Delta t_n}{2}, y_n + \frac{\Delta t_n}{2} k_1\right), \\ y_{n+1} &= y_n + \Delta t_n k_2. \end{aligned}$$

- RK4:

$$\begin{aligned} k_1 &= f(t_n, y_n), \\ k_2 &= f\left(t_n + \frac{\Delta t_n}{2}, y_n + \frac{\Delta t_n}{2} k_1\right), \\ k_3 &= f\left(t_n + \frac{\Delta t_n}{2}, y_n + \frac{\Delta t_n}{2} k_2\right), \\ k_4 &= f(t_n + \Delta t_n, y_n + \Delta t_n k_3), \\ y_{n+1} &= y_n + \frac{\Delta t_n}{6} (k_1 + 2k_2 + 2k_3 + k_4). \end{aligned}$$

Convergence (Thm. 8.10) If the one-step method is stable with respect to perturbations and consistent, then it is convergent. If it is C -stable and of order p , and if f is of class C^p , then

$$\max_{0 \leq n \leq N} |e_n| \leq C(|e_0| + K\Delta t^p) \quad \text{with } K > 0 \text{ indep. of } \Delta t.$$

Stability under perturbations (Thm. 8.11) Assume that Φ is uniformly L -Lipschitz in y with respect to t . Then the one-step method is stable under perturbations with $C = \exp(LT)$.

One-step method in vector form For the test equation $y' = \lambda y$:

$$y_{n+1} = r(\lambda \Delta t_n) y_n.$$

For instance, for RK1 and RK2:

$$r(z) = 1 + z \quad (\text{RK 1}), \quad r(z) = 1 + z + \frac{z^2}{2} \quad (\text{RK 2}).$$

If a Runge–Kutta method is of order p then (Lem. 8.22)

$$r(z) = \exp(z) + C z^{p+1} + O(z^{p+2}).$$

Absolute stability domain (Def. 8.23) It is the set

$$D = \{z \in \overline{\mathbb{C}} : |r(z)| \leq 1\},$$

where $\overline{\mathbb{C}}$ denotes the extended complex plane, i.e., the complex plane with a point at infinity ∞ . The interior is denoted by $\mathring{D}$.

$\mathcal{A}$ -stability (Def. 8.24) A RK method is said to be $\mathcal{A}$ -stable if

$$\{z \in \overline{\mathbb{C}} : \Re(z) \leq 0 \text{ or } z = \infty\} \subset D.$$

Using the maximum principle on $r(z)$, one can show that a method is $\mathcal{A}$ -stable if and only if (Rem. 8.25)

$$\{z \in \mathbb{C} : \Re(z) < 0\} \subset \mathring{D}.$$

Explicit Runge–Kutta methods (Rem. 8.26) Note that an explicit Runge–Kutta method is never $\mathcal{A}$ -stable, since $r(z)$ is a polynomial and its absolute stability domain is necessarily bounded.

Effect of order on stability For RK methods, the stability region grows with order, unlike multistep explicit Adams–Bashforth (AB) methods, whose stability region shrinks as the order increases.

Chapter 9 – PDEs

Heat equation Consider the heat equation:

$$\begin{cases} u_t - \nu u_{xx} = 0, & t \geq 0, x \in [0, 1], \\ u(0, x) = u_0(x), & x \in [0, 1]. \end{cases}$$

To have a unique solution, we must provide boundary conditions. Common choices are Dirichlet (specify u) or Neumann (specify u_x).

Discretization We consider the interval $[0, 1]$ and define

$$\Delta x = 1/(I+1), \quad x_0 = 0, \dots, x_i = i\Delta x, \dots, x_{I+1} = 1,$$

so that we have $I+2$ spatial points. For the time interval $[0, T]$,

$$\Delta t = T/(N+1), \quad t_0 = 0, \dots, t_n = n\Delta t, \dots, t_{N+1} = T,$$

yielding $N+2$ time points.

Finite difference method Approximate $u(t_n, x_i)$ by U_i^n using an explicit Euler in time and second-order centered differences in space:

$$\begin{cases} U_i^0 = u_0(x_i), \\ \frac{U_i^{n+1} - U_i^n}{\Delta t} - \nu \frac{U_{i-1}^n - 2U_i^n + U_{i+1}^n}{\Delta x^2} = 0. \end{cases}$$

Consistency (Def. 9.12) Let ε_i^n denote the local truncation error obtained by inserting the exact $u_i^n := u(t_n, x_i)$ into the scheme:

$$\varepsilon_i^n := \frac{u_i^{n+1} - u_i^n}{\Delta t} - \nu \frac{u_{i-1}^n - 2u_i^n + u_{i+1}^n}{\Delta x^2}.$$

The scheme is said to be consistent if

$$\varepsilon_i^n \rightarrow 0 \quad \text{as } \Delta t \rightarrow 0 \text{ and } \Delta x \rightarrow 0.$$

Order q in time and order p in space if $\varepsilon_i^n = O(\Delta t^q) + O(\Delta x^p)$.

Stability (Def. 9.16) Stable in a norm $\|\cdot\|$ if there exists a constant $K > 0$, independent of Δt and Δx , such that

$$\|U^n\| \leq K \|U^0\|, \quad \forall n \geq 0,$$

for any initial data U^0 . If this holds only under conditions on Δt and Δx , it is said to be conditionally stable.

Stability with matrix norm (Prop. 9.17) Write the discretization scheme as $U^{n+1} = MU^n$. Stability is then equivalent to

$$\|M^n\| \leq \|M\|^n \leq K, \quad n = 0, \dots, N,$$

independently of Δx and Δt (possibly under some conditions).

Convergence (Thm. 9.22) Let $u(t, x)$ be a sufficiently smooth solution of the heat equation with given initial and boundary conditions. Consider a linear two-level scheme (such as above) that is consistent and stable in a norm $\|\cdot\|$. Then the scheme is convergent:

$$\forall T > 0, \quad \lim_{\Delta t \rightarrow 0, \Delta x \rightarrow 0} \sup_{t_n \leq T} \|e^n\| = 0,$$

where $e_j^n = U_j^n - u_j^n$ is the numerical error. Moreover, if the scheme is of order q in time and p in space, then

$$\sup_{t_n \leq T} \|e^n\| = C_T (O(\Delta t^q) + O(\Delta x^p)),$$

for some constant $C_T > 0$ depending on T .

Proving ℓ^2 -stability Expand

$$U^n(x) = \sum_{k \in \mathbb{Z}} \hat{U}^n(k) e^{2i\pi k x},$$

and substitute the Fourier modes into scheme. Using the fact that $U^n(x + \Delta x)$ has coefficients $\hat{U}^n(k) e^{2i\pi k \Delta x}$, we get

$$\hat{U}^{n+1}(k) = a(k) \hat{U}^n(k), \quad a(k) = 1 - 4r \sin^2(\pi k \Delta x).$$

Requiring $|a(k)| \leq 1$ for all k yields

$$\Delta t \leq \frac{\Delta x^2}{2\nu}.$$

Thus the explicit scheme is conditionally ℓ^2 -stable.